


Big models

GPT · Claude · Gemini
input text




100+
layers

next word

~ hundreds of billions of params

The matchbox


TinyStories-1 Layer-21M




1 layer

next word

~ 21 M params



LoRA: a thin patch on top of a frozen model



trainable
 ~ 16 k parameters

64 styles in a matchbox



LoRA
~0.26 % of weights



*children's
stories*

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15

16 attention heads

e.g. Poe

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15

+ LoRA adapter — wakes up the style